

# Lab03 - PySpark

Wei-Lun Tseng (Eric)

ECE, NYCU

[eric840610.ee02@g2.nctu.edu.tw](mailto:eric840610.ee02@g2.nctu.edu.tw)

2022/05/26

# Class Schedule

Total: 3 points

- Lab 3-1: PySpark Warmup
- Lab 3-2: PySpark Exercise

# Checkpoint

## Lab 3-1\_PySpark WarmUp

- Checkpoint 1: Successfully install PySpark on Colab (0.5 point)
- Checkpoint 2: How many people whose age larger than 18? (0.5 point)

## Lab 3-2

- Checkpoint 3: format Open, High, Low, Close, Volume, Adj Close (0.5 point)
- Checkpoint 4: What is the max and min of the Volume column? (0.5 point)
- Checkpoint 5: What is the average Close for each Calendar Month? (0.5 point)

Checkpoint 6: Upload screenshot in correct naming rule (0.5 point)

# Reminder

- The screenshot needs contain coding block and results
- Example:

▼ Use describe() to learn about the DataFrame.

code

```
[8] 1 df.describe().show()
```

summary	Date	Open	High	Low	Close	Volume	Adj Close
count	1258	1258	1258	1258	1258	1258	1258
mean	null	72.35785375357709	72.83938807631165	71.9186009594594	72.38844998012726	8222093.481717011	67.23883848728146
stddev	null	6.76809024470826	6.768186808159218	6.744075756255496	6.756859163732991	4519780.8431556	6.722609449996857
min	2012-01-03	56.389998999999996	57.060001	56.299999	56.419998	2094900	50.363689
max	2016-12-30	90.800003	90.970001	89.25	90.470001	80898100	84.91421600000001

result

# Python and Spark

- In this course the main way we will be working with Python and Spark is through the DataFrame Syntax.
- If you've worked with pandas in Python, R, SQL or even Excel, a DataFrame will feel very familiar!

# Python and Spark

- Spark DataFrames hold data in a column and row format.
- Each column represents some feature or variable.
- Each row represents an individual data point.

# Python and Spark

- Spark began with something known as the “RDD” syntax which was a little ugly and tricky to learn.
- Now Spark 2.0 and higher has shifted towards a DataFrame syntax which is much cleaner and easier to work with!

# Python and Spark

- Spark DataFrames are able to input and output data from a wide variety of sources.
- We can then use these DataFrames to apply various transformations on the data.



# Python and Spark

- At the end of the transformation calls, we can either show or collect the results to display or for some final processing.
- In this section we'll cover all the main features of working with DataFrames that you need to know.

# Python and Spark

- Once we have a solid understanding of Spark DataFrames, we can move on to utilizing the DataFrame MLlib API for Machine Learning.

# Lab 3-1

- PySpark - Warm Up



google colab



全部

影片

圖片

書籍

購物

更多

設定

工具

約有 11,100,000 項結果 (搜尋時間：0.30 秒)

<https://colab.research.google.com> > ... ▼

## Colab - Google

Colaboratory (簡稱為「**Colab**」) 可讓你在瀏覽器上撰寫及執行Python，且具備下列優點：. 不必進行任何設定; 免費使用GPU; 輕鬆共用. 無論你是學生、數據資料學家 ...

<https://research.google.com> > collaboratory ▼ [翻譯這個網頁](#)

## Welcome To Colaboratory - Google Research

With **Colab** you can import an image dataset, train an image classifier on it, and evaluate the model, all in just a few lines of code. **Colab** notebooks execute code ...

[Mounting Google Drive locally](#) · [Google Colab FAQ](#) · [Google Colab](#) · [Colab Widgets](#)

<https://medium.com> > 透過-google-colaboratory-學習... ▼

## 透過Google Colaboratory 學習使用Python 做機器學習等科學 ...

Google Colaboratory ( 以下簡稱**Google Colab** ) 是一個基於Jupyter Notebook 的免費服務 ( 須註冊一個Google 帳號、其餘部份至少撰文的此刻仍是免費 )，所以 ...

<https://www.bnext.com.tw> > article > recommend-to-pr... ▼



歡迎使用 Colaboratory

檔案 編輯 檢視畫面 插入 執行階段 工具

共用



登入

目錄

開始使用

數據資料學

機器學習

其他資源

機器學習範例

區段

說明

常見問題

搜尋程式碼片段

Ctrl+Alt+P

回報錯誤

在 Stack Overflow 上提問

提供意見

查看英文版本

輕鬆共用

?

瀏覽器上撰寫及執行 Python，且具備下列優點：

無論你是學生、數據資料學家或是 AI 研究人員，Colab 都能讓你的工作事半功倍。請觀看 [Colab 的簡介影片](#)瞭解詳情，或是直接瀏覽以下的新手入門說明！

## 開始使用

你正在閱讀的文件並非靜態網頁，而是名為 **Colab 筆記本** 的互動式環境，可讓你撰寫和執行程式碼。

舉例來說，以下是包含簡短 Python 指令碼的**程式碼儲存格**，可進行運算、將值儲存至變數中並列印運算結果：

```
[ ] seconds_in_a_day = 24 * 60 * 60
seconds_in_a_day

86400
```

如要執行上方儲存格中的程式碼，請按一下進行選取，再按一下程式碼左側的播放鍵，或是使用鍵盤快速鍵「Command/Ctrl + Enter 鍵」。按一



## Table of contents

- Getting started
- Data science
- Machine learning
- More Resources
- Machine Learning Examples
- Section

+ Code + Text Copy to Drive

Connect

Editing



## What is Colaboratory?

Colaboratory, or "Colab" for short, allows you to write and execute Python in your browser, with

- Zero configuration required
- Free access to GPUs
- Easy sharing

Whether you're a **student**, a **data scientist** or an **AI researcher**, Colab can make your work easier. Watch [Introduction to Colab](#) to learn more, or just get started below!

## Getting started

The document you are reading is not a static web page, but an interactive environment called a **Colab notebook** that lets you write and execute code.

For example, here is a **code cell** with a short Python script that computes a value, stores it in a variable, and prints the result:

```
[ ] seconds_in_a_day = 24 * 60 * 60
    seconds_in_a_day
```

86400

To execute the code in the above cell, select it with a click and then either press the play button to the left of the code, or use the keyboard shortcut "Command/Ctrl+Enter". To edit the code, just click the cell and start editing.

Login your account...

# Lab 3-1: PySpark Warm Up

- File:
  - 2022-Lab-3-1\_PySpark WarmUp.ipynb
  - people.json
- Checkpoint:
  - Checkpoint 1: Successfully install PySpark on Colab (0.5 pts)
  - Checkpoint 2: How many people whose age smaller than 30? (0.5 pts)





Ta

New notebook

Open notebook

Ctrl+O

Code + Text

Copy to Drive

+ Code + Text

Connect

Editing



## Installation for Pyspark

```
[ ] 1 !apt-get -y install openjdk-8-jre-headless
    2 !pip install pyspark
```

## Start a simple Spark Session

```
[ ] 1 import pyspark
    2 from pyspark.sql import SparkSession
    3 from pyspark.sql.types import StringType, StructType, IntegerType, StructField
    4 spark = SparkSession.builder.appName('Warmup').getOrCreate()
```

### Data Schema

```
[ ] 1 data_schema = [StructField('age', IntegerType(), True),
    2               StructField('name', StringType(), True)]
    3 final_struct = StructType(fields=data_schema)
    4
```

Load the people.json file, have Spark infer the data types.

```
[ ] 1 df = spark.read.json('people.json', schema=final_struct)
```

To execute the code in the above cell, select it with a click and then either press the shortcut "Command/Ctrl+Enter" To edit the code, just click the cell and start editing.

+ Code + Text

Connecting

## Installation for Pyspark

Files

Upload to session storage  
..  
sample\_data

+ Code + Text

RAM  
Disk  
Editing

## Installation for Pyspark

```
[ ] 1 !apt-get -y install openjdk-8-jre-headless  
    2 !pip install pyspark
```

### Start a simple Spark Session

```
[ ] 1 import pyspark  
    2 from pyspark.sql import SparkSession  
    3 from pyspark.sql.types import StringType, StructType, IntegerType, StructField  
    4 spark = SparkSession.builder.appName('Warmup').getOrCreate()
```

Data Schema

## ▼ Installation for Pyspark

# Checkpoint-1: Successfully installed Pyspark 0.5 point

```
[1] 1 !apt-get -y install openjdk-8-jre-headless
    2 !pip install pyspark
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
Suggested packages:
  libnss-mdns fonts-dejavu-extra fonts-ipafont-gothic fonts-ipafont-mincho
  fonts-wqy-microhei fonts-wqy-zenhei fonts-indic
The following NEW packages will be installed:
  openjdk-8-jre-headless
0 upgraded, 1 newly installed, 0 to remove and 39 not upgraded.
Need to get 28.2 MB of archives.
After this operation, 104 MB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 openjdk-8-jre-headless amd64 8u292-b10-0ubuntu1~18.04 [28.2 MB]
Fetched 28.2 MB in 2s (15.9 MB/s)
Selecting previously unselected package openjdk-8-jre-headless:amd64.
(Reading database ... 160772 files and directories currently installed.)
Preparing to unpack .../openjdk-8-jre-headless_8u292-b10-0ubuntu1~18.04_amd64.deb ...
Unpacking openjdk-8-jre-headless:amd64 (8u292-b10-0ubuntu1~18.04) ...
Setting up openjdk-8-jre-headless:amd64 (8u292-b10-0ubuntu1~18.04) ...
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/orbd to provide /usr/bin/orbd (orbd) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/servertool to provide /usr/bin/servertool (servertool) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/tnameserv to provide /usr/bin/tnameserv (tnameserv) in auto mode
Collecting pyspark
  Downloading https://files.pythonhosted.org/packages/89/db/e18cfd78e408de957821ec5ca56de1250645b05f8523d169803d8df35a64/pyspark-3.1.2.tar.gz (212.4MB)
    |████████████████████████████████████████| 212.4MB 63kB/s
Collecting py4j==0.10.9
  Downloading https://files.pythonhosted.org/packages/9e/b6/6a4fb90cd235dc8e265a6a2067f2a2c99f0d91787f06aca4bcf7c23f3f80/py4j-0.10.9-py2.py3-none-any.whl
    |████████████████████████████████████████| 204kB 19.1MB/s
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.1.2-py2.py3-none-any.whl size=212880768 sha256=b63ae7fb090bc38cdd82055abe3e77e4b9b8aee1d84c4cb3a9b36193
  Stored in directory: /root/.cache/pip/wheels/40/1b/2c/30f43be2627857ab80062bef1527c0128f7b4070b6b2d02139
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9 pyspark-3.1.2
```

## ▼ Start a simple Spark Session

```
[2] 1 import pyspark
    2 from pyspark.sql import SparkSession
    3 from pyspark.sql.types import StringType, StructType, IntegerType, StructField
    4 spark = SparkSession.builder.appName('Warmup').getOrCreate()
```

### Data Schema

```
[3] 1 data_schema = [StructField('age', IntegerType(), True),
    2                 StructField('name', StringType(), True)]
    3 final_struct = StructType(fields=data_schema)
    4
```

Load the people.json file, have Spark infer the data types.

```
[5] 1 df = spark.read.json('people.json', schema=final_struct)
```

## ▼ What are the column names?

```
[6] 1 df.columns

['age', 'name']
```

## ▼ What is the schema?

```
[7] 1 df.printSchema()

root
 |-- age: integer (nullable = true)
 |-- name: string (nullable = true)
```

### Show whole DataFrame

```
1 df.show()
2
```

```
+---+-----+
| age | name |
+---+-----+
| null | Michael |
| 30 | Andy |
| 19 | Justin |
+---+-----+
```

Print out the first 2 rows.

```
[9] 1 # Didn't strictly need a for loop, could have just then head()
    2 for row in df.head(2):
    3     print(row)
    4     print('\n')
```

```
Row(age=None, name='Michael')
```

```
Row(age=30, name='Andy')
```

Use describe() to learn about the DataFrame

```
[10] 1 df.describe()
```

```
DataFrame[summary: string, age: string, name: string]
```

Use another data frame to learn about the statistical report

```
[11] 1 temp = df.describe()
    2 temp.show()
```

```
+-----+-----+-----+
|summary|      age|  name|
+-----+-----+-----+
|  count|         2|      3|
|   mean|      24.5| null|
| stddev|7.7781745930520225| null|
|    min|         19|  Andy|
|    max|         30|Michael|
+-----+-----+-----+
```

There are too many decimal places for mean and stddev in the describe() dataframe.

How to deal with it?

```
[13] 1 from pyspark.sql.functions import format_number
```

```
1 result = df.describe()
2 result.select(result['summary'],
3               format_number(result['age'].cast('float'),2).alias('age'))
4         ).show()
```

```
+-----+-----+
|summary|  age|
+-----+-----+
|  count|  2.00|
|   mean| 24.50|
| stddev|  7.78|
|    min|19.00|
```

Get the mean of age directly

```
[15] 1 from pyspark.sql.functions import mean
      2 df.select(mean("age")).show()
```

```
+-----+
|avg(age)|
+-----+
|    24.5|
+-----+
```

What is the max and min of the Volume column?

```
[16] 1 # Could have also used describe
      2 from pyspark.sql.functions import max,min
      3 df.select(max("age"),min("age")).show()
```

```
+-----+-----+
|max(age)|min(age)|
+-----+-----+
|      30|       19|
+-----+-----+
```

How many days was the age smaller than 30?

```
[17] 1 df.filter("age < 30").count()
```

1

```
1 from pyspark.sql.functions import count
2 result = df.filter(df['age'] < 30)
3 result.select(count('age')).show()
```

```
+-----+
|count(age)|
+-----+
|          1|
+-----+
```

people.json X

```
1 {"name":"Michael"}
2 {"name":"Andy", "age":30}
3 {"name":"Justin", "age":19}
4
```

# Check-Point-2:

0.5 point

**Checkpoint 2 - 0.5 point**


How many people whose age larger than 18?

[ ]

# Lab 3-2

- PySpark - Exercise





+ Code + Text

✓ RAM  
Disk

Editing



## ▼ Spark DataFrames Project Exercise

Let's get some quick practice with your new Spark DataFrame skills, you will be asked some basic questions about some stock market data, in this case Walmart Stock from the years 2012-2017. This exercise will just ask a bunch of questions, unlike the future machine learning exercises, which will be a little looser and be in the form of "Consulting Projects", but more on that later!

For now, just answer the questions and complete the tasks below.

[Hint](#)

- ▼ Use the `walmart_stock.csv` file to Answer and complete the tasks below!

```
[ ] 1 !apt-get -y install openjdk-8-jre-headless
     2 !pip install pyspark
     3
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
openjdk-8-jre-headless is already the newest version (8u275-b01-0ubuntu1~18.04).
0 upgraded, 0 newly installed, 0 to remove and 14 not upgraded.
Requirement already satisfied: pyspark in /usr/local/lib/python3.6/dist-packages (3.0.1)
Requirement already satisfied: py4j==0.10.9 in /usr/local/lib/python3.6/dist-packages (from pyspark) (0.10.9)
```

- ▼ Use the walmart\_stock.csv file to answer and complete the tasks below!

```
[1] 1 !apt-get -y install openjdk-8-jre-headless
    2 !pip install pyspark
    3

Reading package lists... Done
Building dependency tree
Reading state information... Done
Suggested packages:
  libnss-mdns fonts-dejavu-extra fonts-ipafont-gothic fonts-ipafont-mincho
  fonts-wqy-microhei fonts-wqy-zenhei fonts-indic
The following NEW packages will be installed:
  openjdk-8-jre-headless
0 upgraded, 1 newly installed, 0 to remove and 39 not upgraded.
Need to get 28.2 MB of archives.
After this operation, 104 MB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 openjdk-8-jre-headless amd64 8u292-b10-0ubuntu1~18.04 [28.2 MB]
Fetched 28.2 MB in 1s (50.6 MB/s)
Selecting previously unselected package openjdk-8-jre-headless:amd64.
(Reading database ... 160772 files and directories currently installed.)
Preparing to unpack .../openjdk-8-jre-headless_8u292-b10-0ubuntu1~18.04_amd64.deb ...
Unpacking openjdk-8-jre-headless:amd64 (8u292-b10-0ubuntu1~18.04) ...
Setting up openjdk-8-jre-headless:amd64 (8u292-b10-0ubuntu1~18.04) ...
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/orbd to provide /usr/bin/orbd (orbd) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/servertool to provide /usr/bin/servertool (servertool) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/tnameserv to provide /usr/bin/tnameserv (tnameserv) in auto mode
Collecting pyspark
  Downloading https://files.pythonhosted.org/packages/89/db/e18cf78e408de957821ec5ca56de1250645b05f8523d169803d8df35a64/pyspark-3.1.2.tar.gz (212.4MB)
    |████████████████████████████████████████| 212.4MB 71kB/s
Collecting py4j==0.10.9
  Downloading https://files.pythonhosted.org/packages/9e/b6/6a4fb90cd235dc8e265a6a2067f2a2c99f0d91787f06aca4bcf7c23f3f80/py4j-0.10.9-py2.py3-none-any.w
    |████████████████████████████████████████| 204kB 21.1MB/s
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.1.2-py2.py3-none-any.whl size=212880768 sha256=9dbb7b7ed3d5afe50b472af777003fe07c00f5516b6d6bedadc985bc
  Stored in directory: /root/.cache/pip/wheels/40/1b/2c/30f43be2627857ab80062bef1527c0128f7b4070b6b2d02139
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9 pyspark-3.1.2
```

- ▼ Start a simple Spark Session

```
[2] 1 from pyspark.sql import SparkSession
    2 spark = SparkSession.builder.appName("walmart").getOrCreate()
```

- ▼ Load the Walmart Stock CSV File, have Spark infer the data types.

```
[3] 1 df = spark.read.csv('walmart_stock.csv',header=True,inferSchema=True)
```

▼ What are the column names?

```
[4] 1 df.columns  
  
['Date', 'Open', 'High', 'Low', 'Close', 'Volume', 'Adj Close']
```

▼ What does the Schema look like?

```
[6] 1 df.printSchema()  
  
root  
|-- Date: string (nullable = true)  
|-- Open: double (nullable = true)  
|-- High: double (nullable = true)  
|-- Low: double (nullable = true)  
|-- Close: double (nullable = true)  
|-- Volume: integer (nullable = true)  
|-- Adj Close: double (nullable = true)
```

▼ Print out the first 5 columns.

```
[7] 1 # Didn't strictly need a for loop, could have just then head()  
2 for row in df.head(5):  
3     print(row)  
4     print('\n')
```

Row(Date='2012-01-03', Open=59.970001, High=61.060001, Low=59.869999, Close=60.330002, Volume=12668800, Adj Close=52.619234999999996)

Row(Date='2012-01-04', Open=60.209998999999996, High=60.349998, Low=59.470001, Close=59.709998999999996, Volume=9593300, Adj Close=52.078475)

Row(Date='2012-01-05', Open=59.349998, High=59.619999, Low=58.369999, Close=59.419998, Volume=12768200, Adj Close=51.825539)

Row(Date='2012-01-06', Open=59.419998, High=59.450001, Low=58.869999, Close=59.0, Volume=8069400, Adj Close=51.45922)

Row(Date='2012-01-09', Open=59.029999, High=59.549999, Low=58.919998, Close=59.18, Volume=6679300, Adj Close=51.616215000000004)

▼ Use describe() to learn about the DataFrame.

```
[8] 1 df.describe().show()
```

summary	Date	Open	High	Low	Close	Volume	Adj Close
count	1258	1258	1258	1258	1258	1258	1258
mean	null	72.35785375357709	72.83938807631165	71.9186009594594	72.38844998012726	8222093.481717011	67.23883848728146
stddev	null	6.76809024470826	6.768186808159218	6.744075756255496	6.756859163732991	4519780.8431556	6.722609449996857
min	2012-01-03	56.389998999999996	57.060001	56.299999	56.419998	2094900	50.363689
max	2016-12-30	90.800003	90.970001	89.25	90.470001	80898100	84.91421600000001

```
[9] 1 # Uh oh Strings!
    2 df.describe().printSchema()
```

```
root
|-- summary: string (nullable = true)
|-- Date: string (nullable = true)
|-- Open: string (nullable = true)
|-- High: string (nullable = true)
|-- Low: string (nullable = true)
|-- Close: string (nullable = true)
|-- Volume: string (nullable = true)
|-- Adj Close: string (nullable = true)
```

```
[12] 1 # hint
     2 from pyspark.sql.functions import format_number
```



```
1 result = df.describe()
2 result.select(result['summary'],
3               format_number(result['Open'].cast('float'),2).alias('Open'),
4               result['Volume'].cast('int').alias('Volume'))
5 .show()
```



```
+-----+-----+-----+
|summary|    Open|  Volume|
+-----+-----+-----+
|  count|1,258.00|   1258|
|   mean|   72.36| 8222093|
| stddev|    6.77| 4519780|
|    min|   56.39| 2094900|
|    max|   90.80| 80898100|
+-----+-----+-----+
```

# Check-Point-3:

## Check Point 3

format Open, High, Low, Close, Volume, Adj Close

0.5 point

[14] 1

summary	Open	High	Low	Close	Volume
count	1,258.00	1,258.00	1,258.00	1,258.00	1258
mean	72.36	72.84	71.92	72.39	8222093
stddev	6.77	6.77	6.74	6.76	4519780
min	56.39	57.06	56.30	56.42	2094900
max	90.80	90.97	89.25	90.47	80898100

▼ Create a new dataframe with a column called HV Ratio that is the ratio of the High Price versus volume of stock traded for a day.

```
[15] 1 df2 = df.withColumn("HV Ratio",df["High"]/df["Volume"])#.show()  
2 # df2.show()  
3 df2.select('HV Ratio').show()
```

```
+-----+  
|          HV Ratio|  
+-----+  
|4.819714653321546E-6|  
|6.290848613094555E-6|  
|4.669412994783916E-6|  
|7.367338463826307E-6|  
|8.915604778943901E-6|  
|8.644477436914568E-6|  
|9.351828421515645E-6|  
| 8.29141562102703E-6|  
|7.712212102001476E-6|  
|7.071764823529412E-6|  
|1.015495466386981E-5|  
|6.576354146362592...|  
| 5.90145296180676E-6|  
|8.547679455011844E-6|  
|8.420709512685392E-6|  
|1.041448341728929...|  
|8.316075414862431E-6|  
|9.721183814992126E-6|  
|8.029436027707578E-6|  
|6.307432259386365E-6|  
+-----+  
only showing top 20 rows
```



▼ What day had the Peak High in Price?

```
[16] 1 # Didn't need to really do this much indexing
      2 # Could have just shown the entire row
      3 df.orderBy(df["High"].desc()).head(1)[0][0]
```

'2015-01-13'

▼ What is the mean of the Close column?



```
1 # Also could have gotten this from describe()
2 from pyspark.sql.functions import mean # hint
3 df.select(mean("Close")).show()
```

```
+-----+
|      avg(Close)|
+-----+
|72.38844998012726|
+-----+
```





# Check-Point-4:

1 point

## Check Point 4

What is the max and min of the Volume column?

```
[19]  1 # Could have also used describe  
      2 from pyspark.sql.functions import max,min
```

```
[20]  1
```

max(Volume)	min(Volume)
80898100	2094900

▼ How many days was the Close lower than 60 dollars?

```
[24] 1 df.filter("Close < 60").count()
```

```
81
```

```
[23] 1 df.filter(df['Close'] < 60).count()
```

```
81
```

```
[26] 1 from pyspark.sql.functions import count
      2 result = df.filter(df['Close'] < 60)
      3 result.select(count('Close')).show()
```

```
+-----+
|count(Close)|
+-----+
|          81|
+-----+
```

▼ What percentage of the time was the High greater than 80 dollars ?

In other words, (Number of Days High>80)/(Total Days in the dataset)

```
[27] 1 # Many ways to do this
      2 (df.filter(df["High"]>80).count()*1.0/df.count())*100
```

```
9.141494435612083
```

▼ What is the Pearson correlation between High and Volume?

```
[28] 1 from pyspark.sql.functions import corr
      2 df.select(corr("High","Volume")).show()
```

```
+-----+
|corr(High, Volume)|
+-----+
|-0.3384326061737161|
+-----+
```

▼ What is the max High per year?

```
[29] 1 from pyspark.sql.functions import year
      2 year_df = df.withColumn("Year",year(df["Date"]))
```

```
[30] 1 max_df = year_df.groupBy('Year').max()
      2 # 2015
      3 max_df.select('Year','max(High)').show()
```

```
+----+-----+
|Year|max(High)|
+----+-----+
|2015|90.970001|
|2013|81.370003|
|2014|88.089996|
|2012|77.599998|
|2016|75.190002|
+----+-----+
```

▼ What is the max Close for each Calendar Month?

In other words, across all the years, what is the max Close price for Jan, Feb, Mar, etc... Your result will have a value for each of these months.

```
1 from pyspark.sql.functions import month # hint
2 monthdf = df.withColumn("Month",month("Date"))
3 monthmax = monthdf.select("Month","Close").groupBy("Month").max()
4 monthmax.select("Month","max(Close)").orderBy('Month').show()
5
```

Month	max(Close)
1	90.470001
2	87.33000200000001
3	83.959999
4	81.029999
5	79.91999799999999
6	77.32
7	78.550003
8	78.769997
9	77.510002
10	78.290001
11	87.540001
12	86.910004

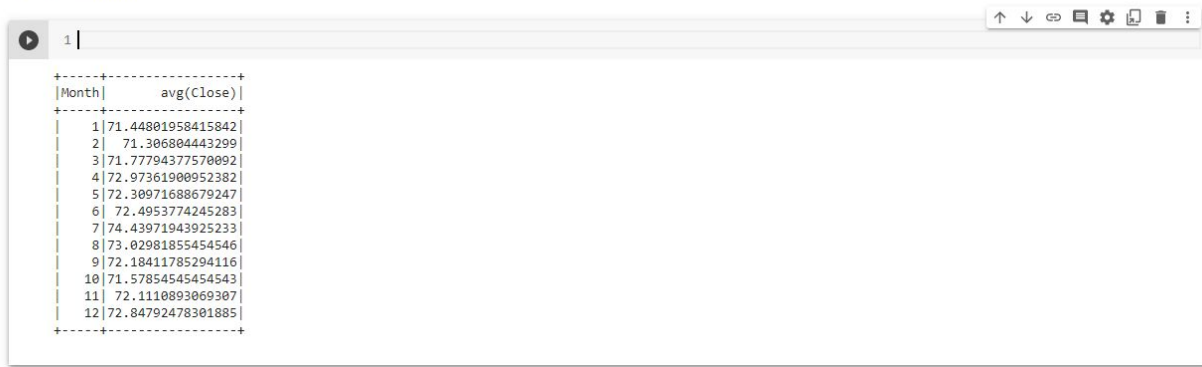
# Check-Point-5:

1 point

## Check Point 5

- ▼ What is the average Close for each Calendar Month?

In other words, across all the years, what is the average Close price for Jan, Feb, Mar, etc... Your result will have a value for each of these months.



The screenshot shows a SQL query result window with a toolbar at the top containing icons for navigation and editing. The query is numbered '1'. The result is a table with two columns: 'Month' and 'avg(Close)'. The data is presented in a grid with dashed borders around the header and each row.

Month	avg(Close)
1	71.44801958415842
2	71.306804443299
3	71.77794377570092
4	72.97361900952382
5	72.30971688679247
6	72.4953774245283
7	74.43971943925233
8	73.02981855454546
9	72.18411785294116
10	71.57854545454543
11	72.1110893069307
12	72.84792478301885

Great ~ !

## Checkpoint 6

Download your jupyter notebook and  
upload to new E3 (Lab.3 PySpark)  
(0.5 points)

# Checkpoint 6

Upload screenshots of 5 checkpoint results to E3.

- File format: jpg/png
- Number of files: 5
- File name:
  - [student-ID]\_ckpt\_[no.].jpg
  - Example: 088888\_ckpt\_1.jpg
- Note: Do not use any folder or zip file to pack your pictures.
  - Ex.

檔案大小上限：2GB，最多附件：20

<input type="checkbox"/>	名稱	最後修改	大小	類型
<input type="checkbox"/>	088888_ckpt_1.png	2022/05/24 15:55	515 位元組	圖像(PNG)
<input type="checkbox"/>	088888_ckpt_2.png	2022/05/24 15:55	515 位元組	圖像(PNG)
<input type="checkbox"/>	088888_ckpt_3.png	2022/05/24 15:56	515 位元組	圖像(PNG)
<input type="checkbox"/>	088888_ckpt_4.png	2022/05/24 15:56	515 位元組	圖像(PNG)
<input type="checkbox"/>	088888_ckpt_5.png	2022/05/24 15:56	515 位元組	圖像(PNG)

## For homework 4:

### PySpark ML lib warmup

#### Dataset Overview



```
1 import pandas as pd
2 df = pd.read_csv('public.csv')
3 df
```



	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	15565701	Ferri	698	Spain	Female	39	9	161993.89	1	0	0	90212.38	0
1	15565706	Akobundu	612	Spain	Male	35	1	0.00	1	1	1	83256.26	1
2	15565796	Docherty	745	Germany	Male	48	10	96048.55	1	1	0	74510.65	0
3	15565806	Toosey	532	France	Male	38	9	0.00	2	0	0	30583.95	0
4	15565878	Bates	631	Spain	Male	29	3	0.00	2	1	1	197963.46	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
7995	15815628	Moysey	711	France	Female	37	8	113899.92	1	0	0	80215.20	0
7996	15815645	Akhtar	481	France	Male	37	8	152303.66	2	1	1	175082.20	0
7997	15815656	Hopkins	541	Germany	Female	39	9	100116.67	1	1	1	199808.10	1
7998	15815660	Mazzi	758	France	Female	34	1	154139.45	1	1	1	60728.89	0
7999	15815690	Akabueze	614	Spain	Female	40	3	113348.50	1	1	1	77789.01	0

8000 rows x 13 columns